



Deep Learning Topics

A literature overview

Yinan Yu
2019-10-08 MLFP

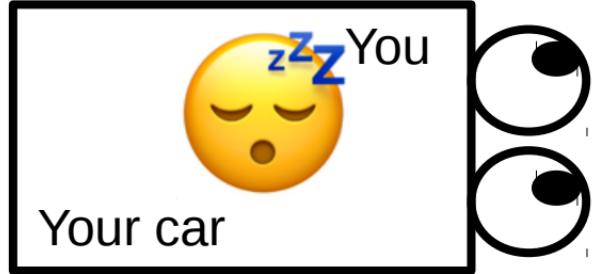


Deep learning!

WHAT FOR???

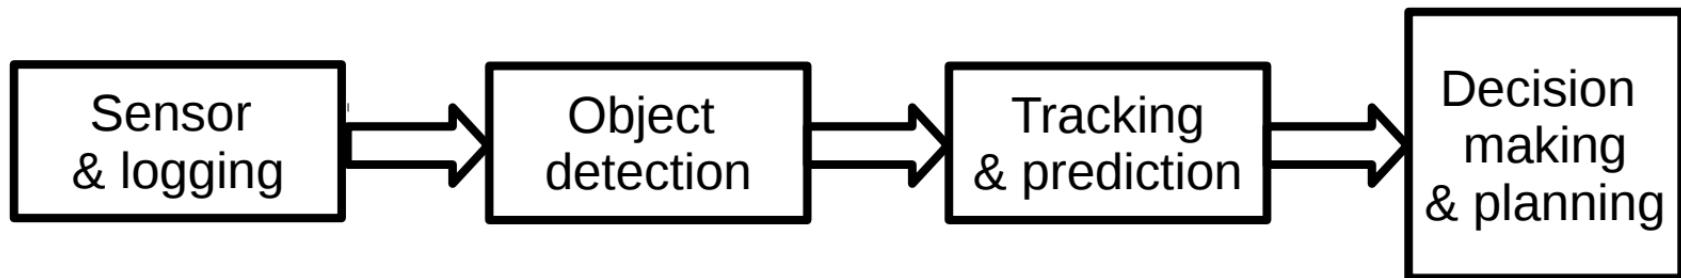


Self-driving cars



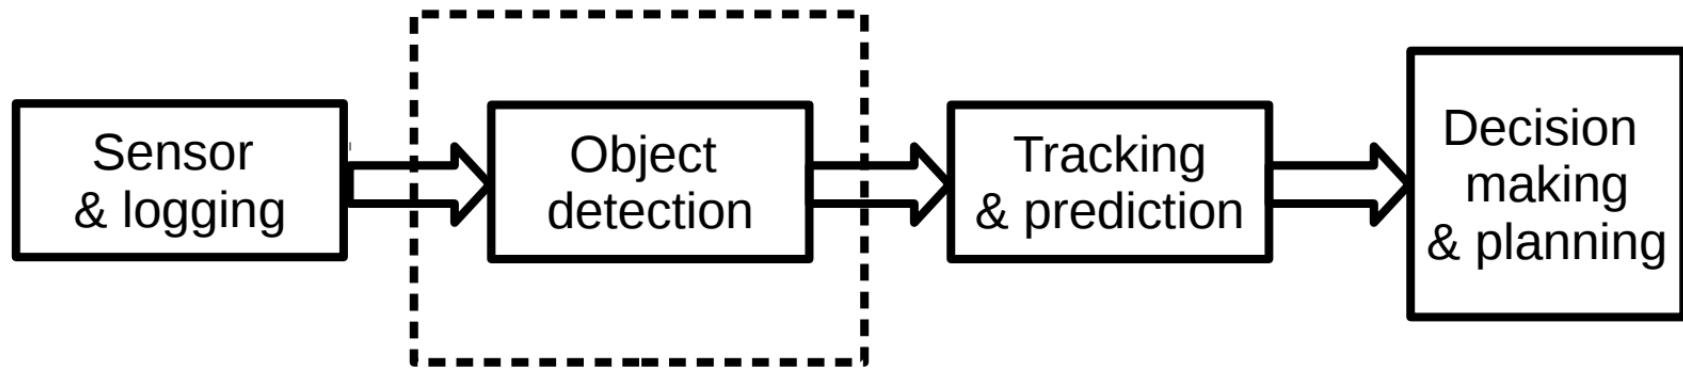


Self-driving cars





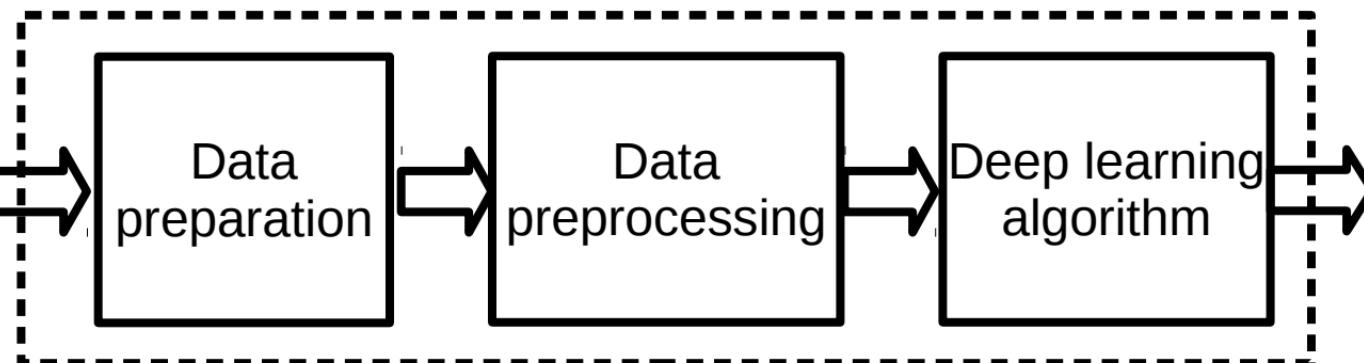
Self-driving cars



Self-driving cars

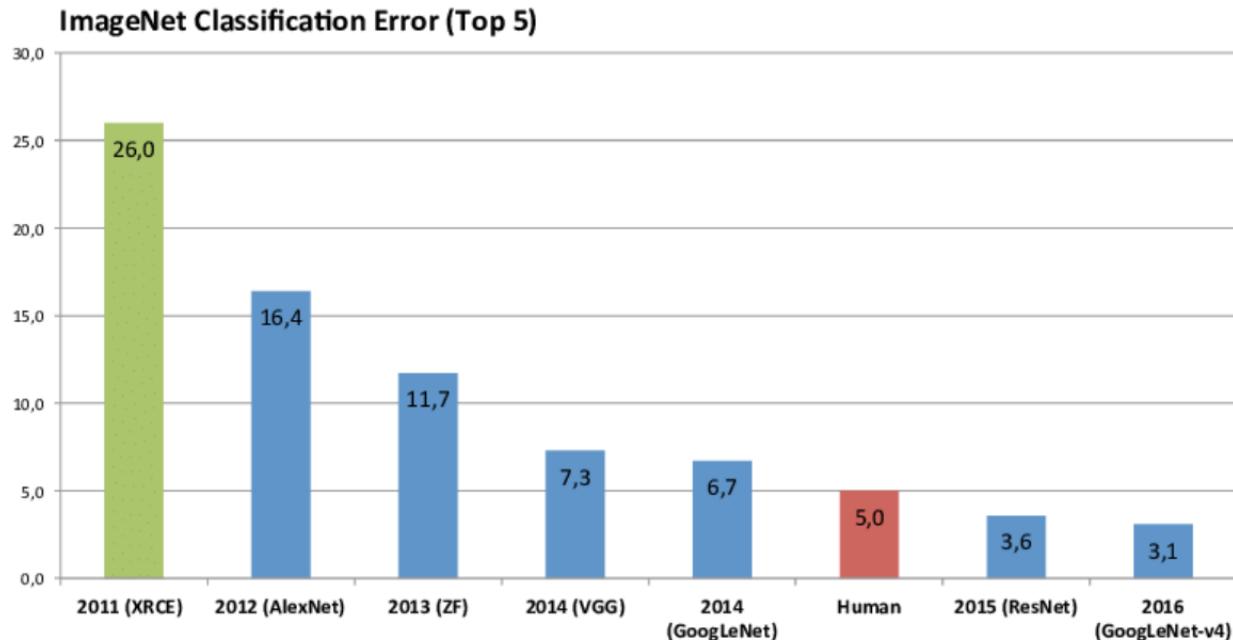


Deep learning





Computer vision (ImageNet) challenge





Purpose

- What are the different building blocks?
- What are the “variations”?
- How are they evaluated?
- Interesting subjects for future MLFP talks?



Outline

- Overview
 - Deep learning pipeline
 - Evaluation: accuracy, complexity, training time, inference time, memory footprint
- Components
 - Data modeling
 - Architecture
 - Loss, optimization
 - Abstraction and annotation
 - Infrastructure

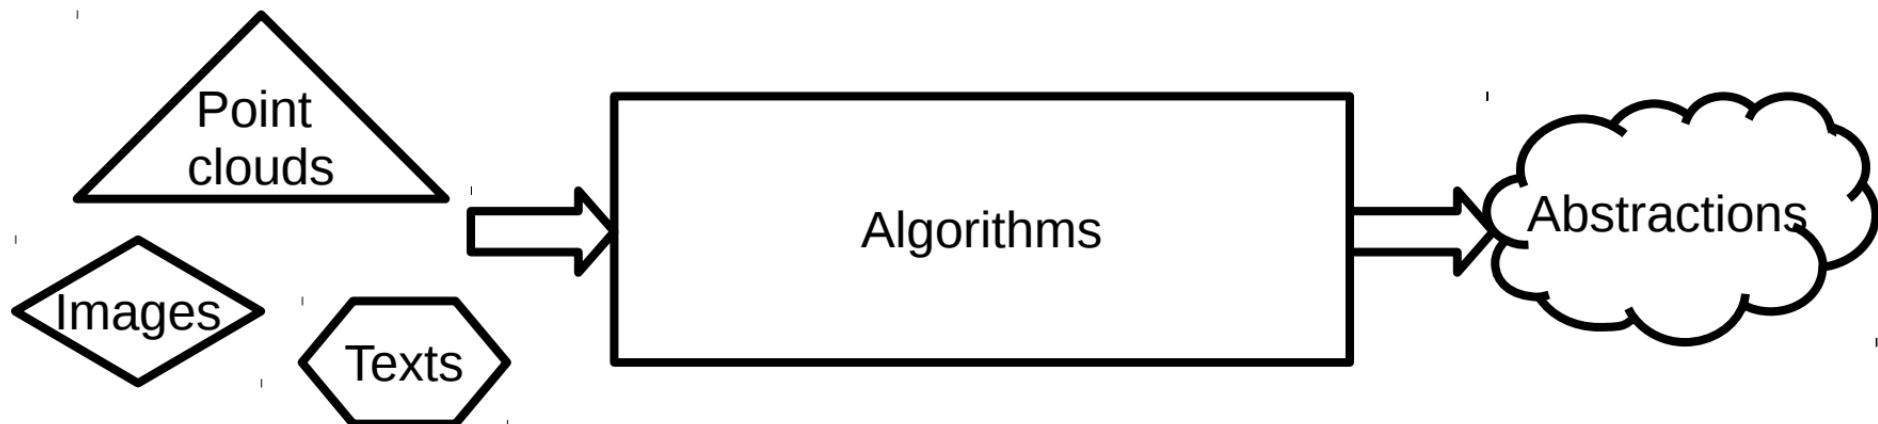
MLFP

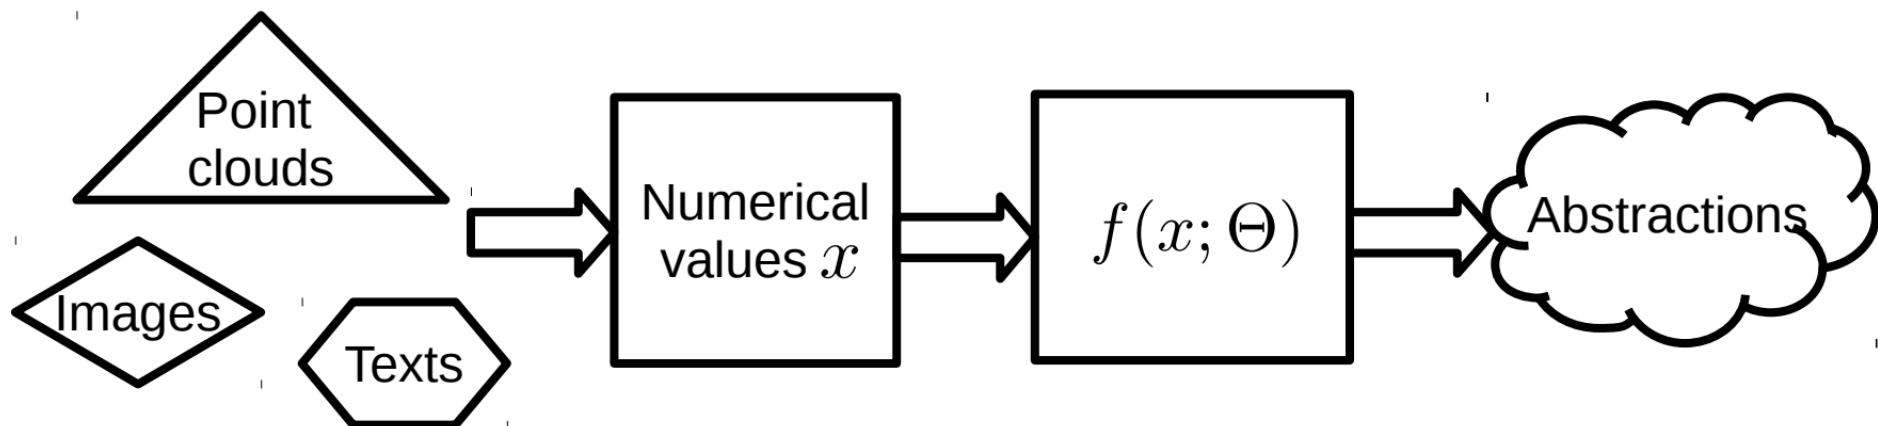


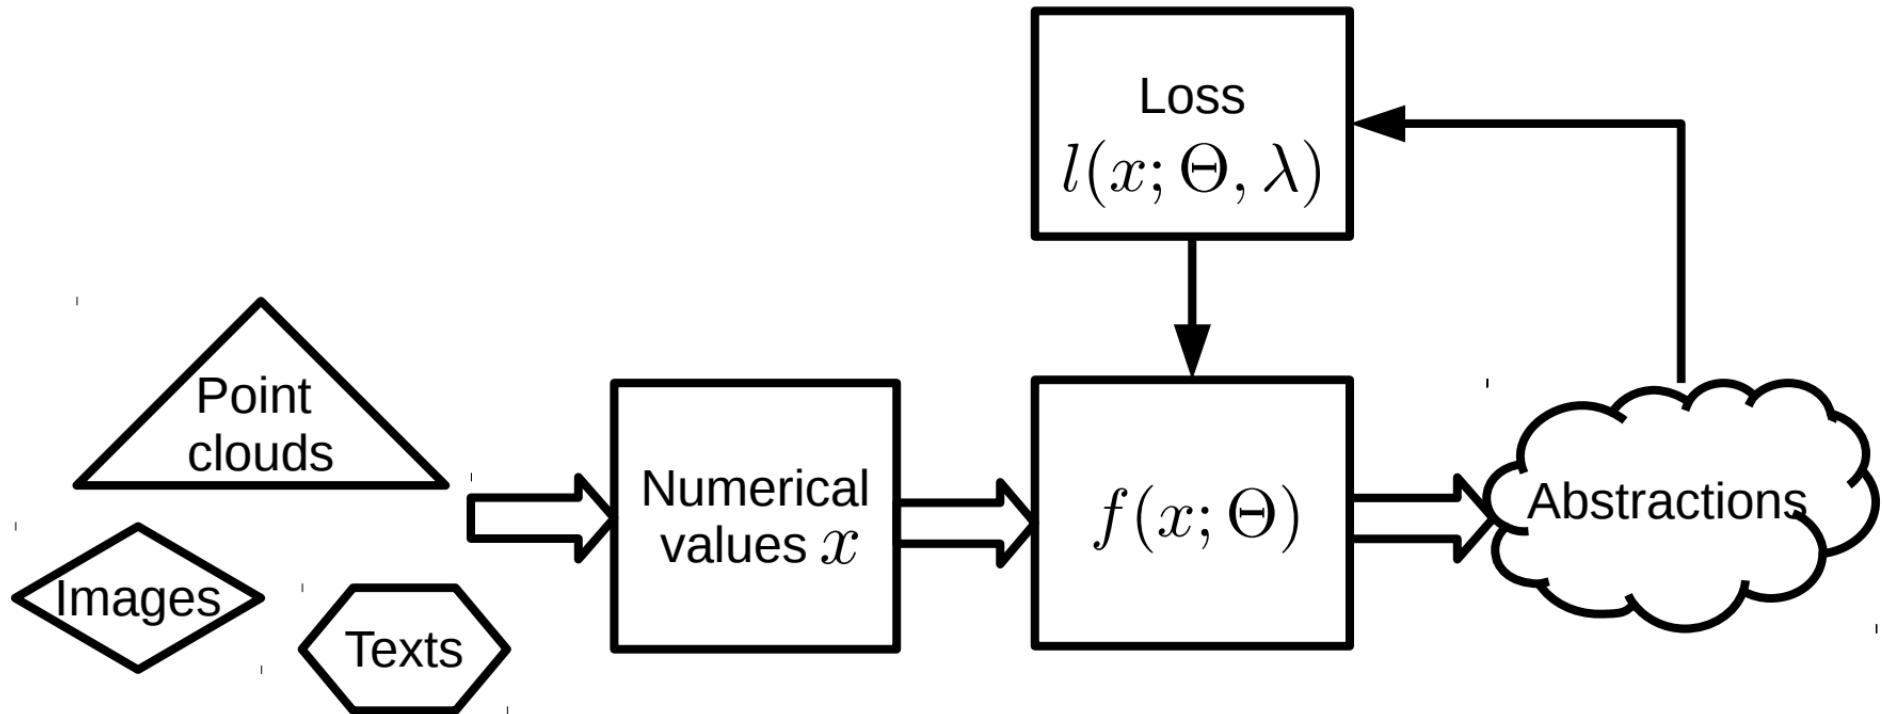
Overview

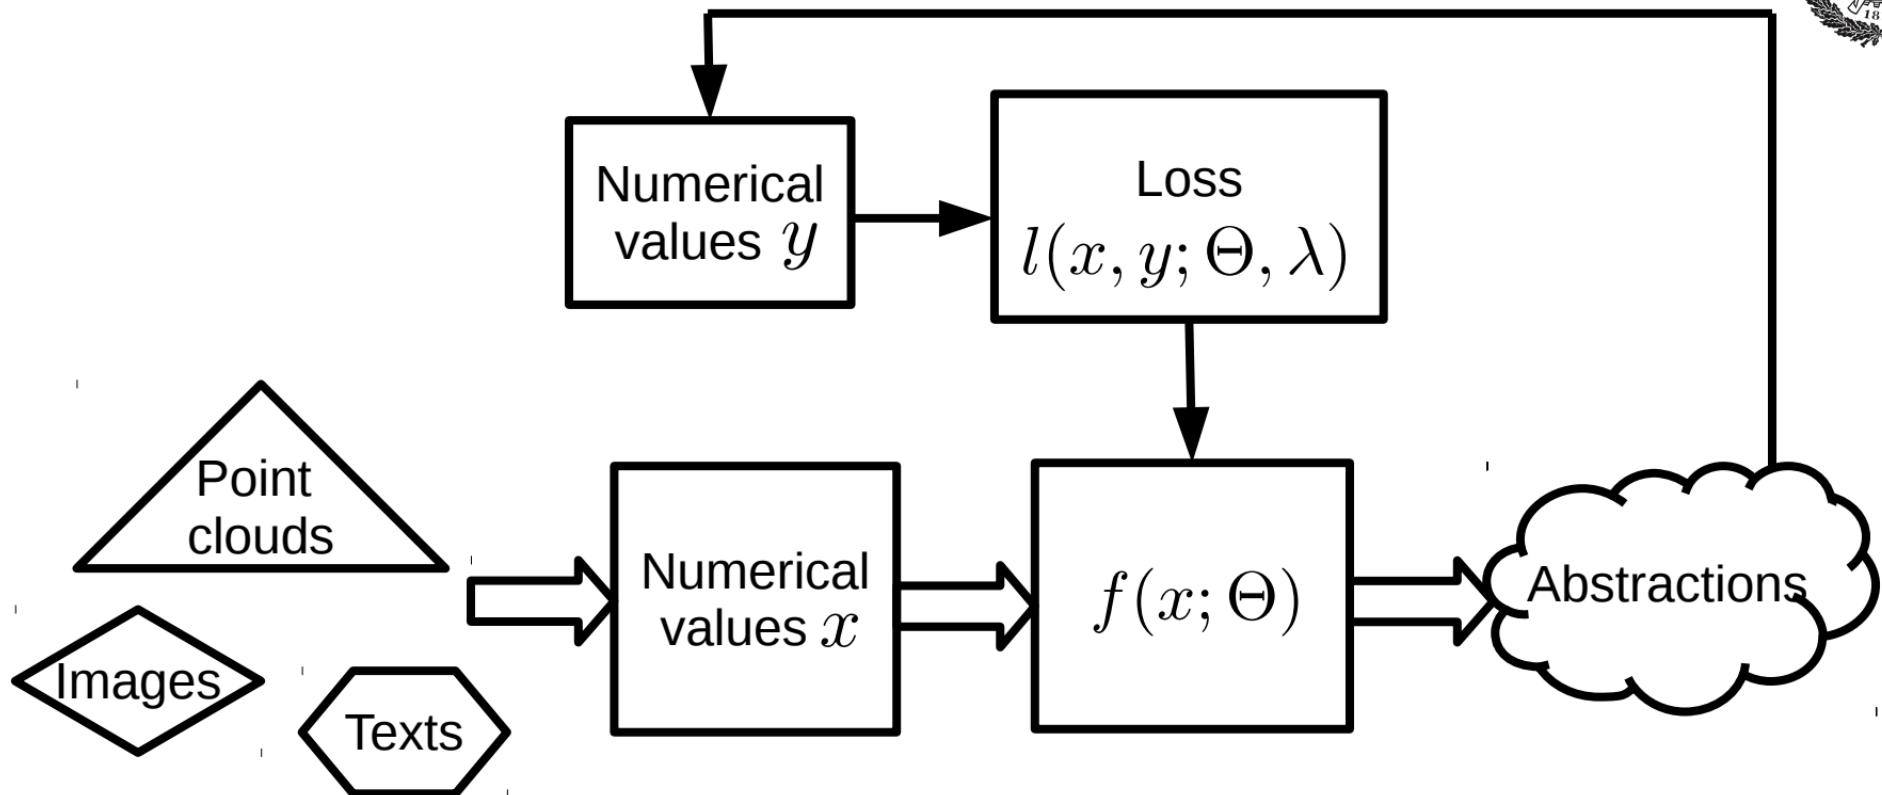


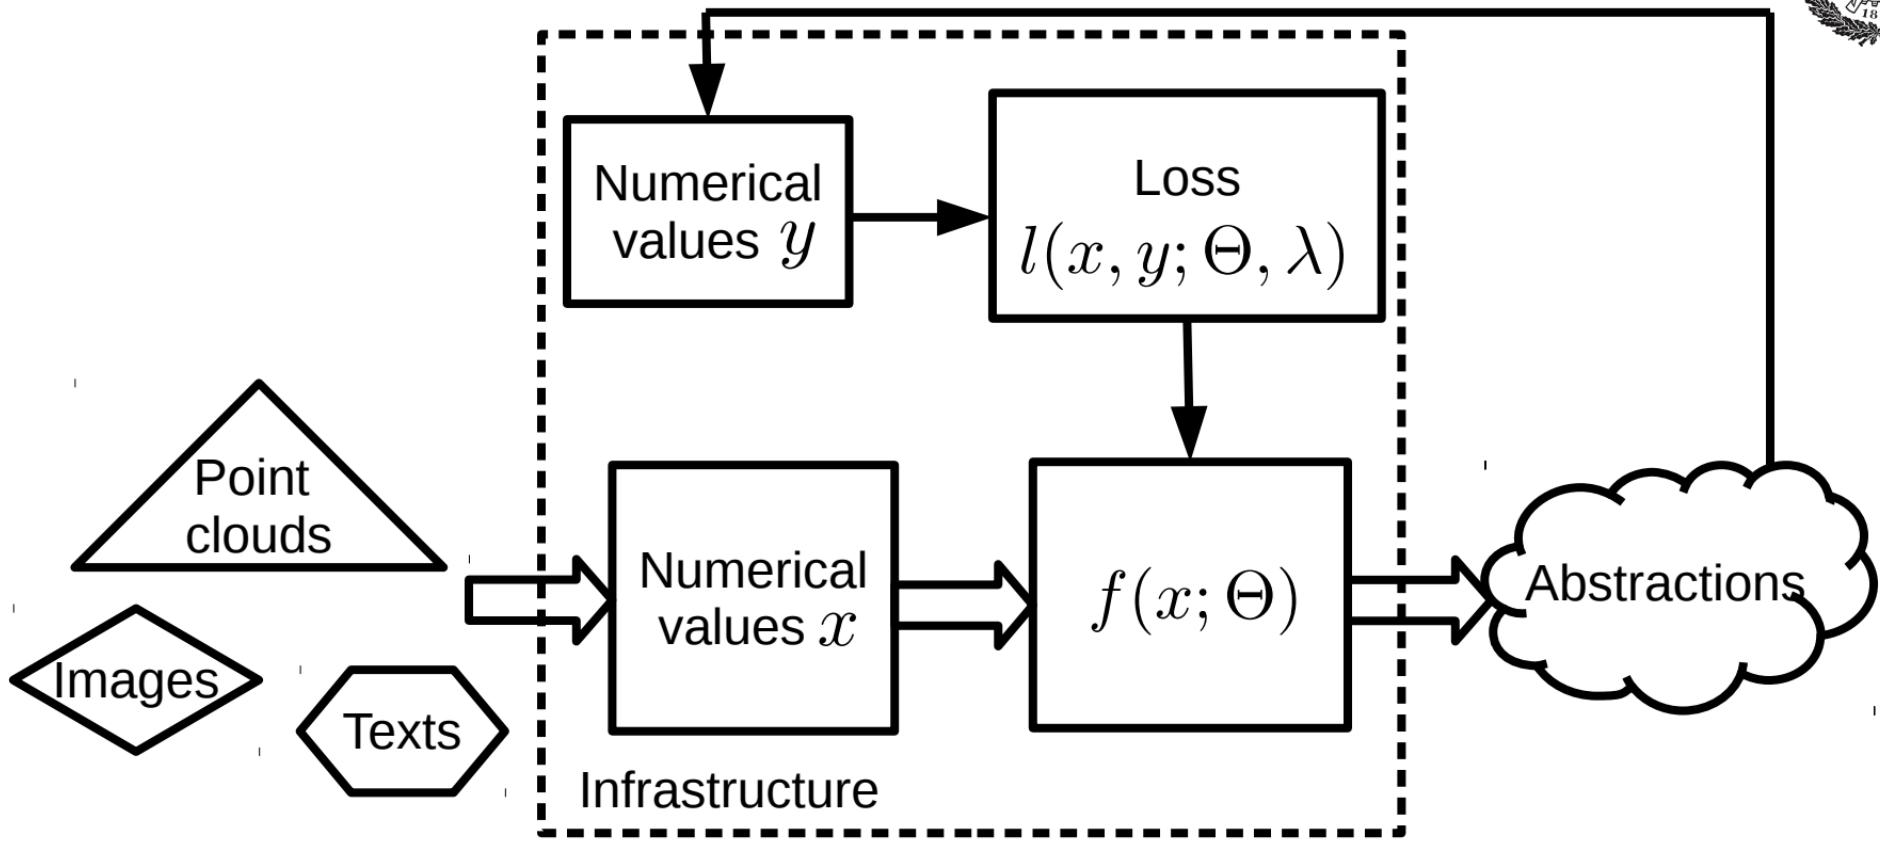
Deep learning: data driven approach













Evaluation

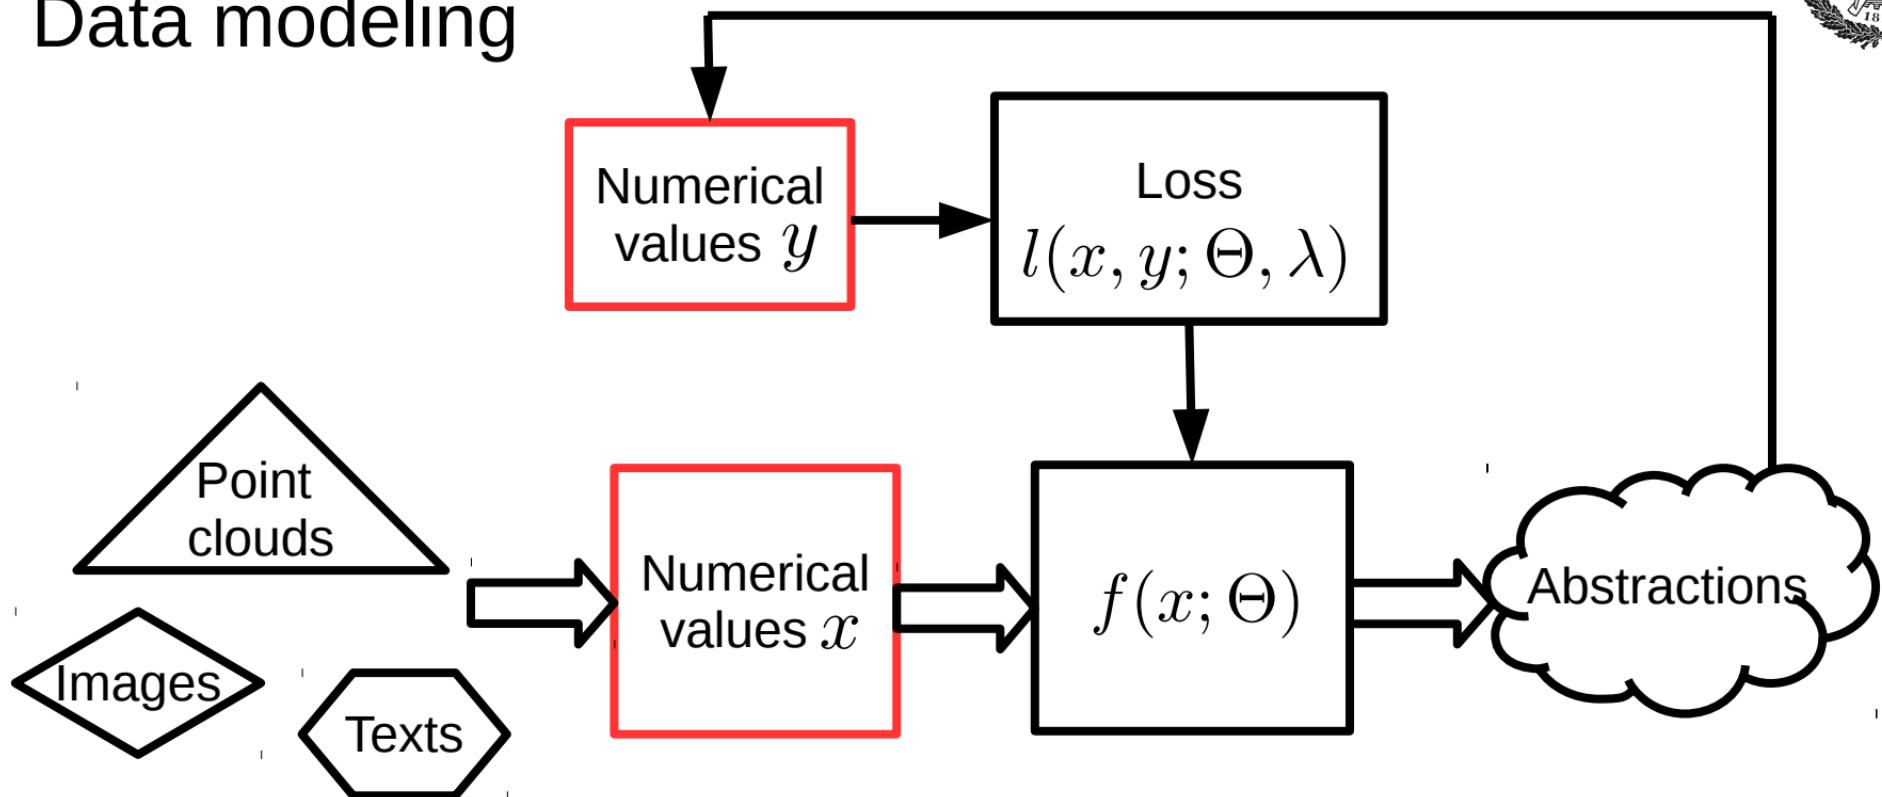
- Accuracy
 - Quantitative error: classification error, regression error
 - Qualitative error: visualization, human interpretation
 - Performance in the “big picture”
- Complexity measures
 - Number of trainable parameters
 - Model specific complexity
- Training time
 - Theoretical time complexity (number of training data, dimensionality, network architecture, error tolerance)
 - Measured training time for specific infrastructures
- Inference time
 - FLOPs (MACC)
 - Measured inference time for specific infrastructures/target hardware
- Memory footprint



Components



Data modeling





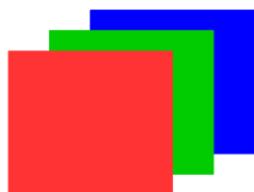
Data modeling

- Goal: transform data (texts, images, point clouds, etc) and their annotations into “meaningful” numerical values
- Challenges:
 - Numerical interpretation from arbitrary data format
 - Meaningful representation of annotations
 - Variable size
 - Unclear contribution to the machine learning task



Data modeling

- Images:



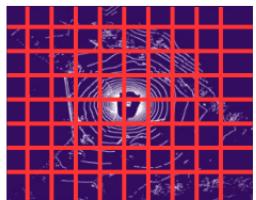
- Text:

- ["cat", "cute", "be", "I", "a", "have"]
- "Cats are cute." → [1, 1, 1, 0, 0, 0]
- "I have a cute cat." → [1, 1, 0, 1, 1, 1]

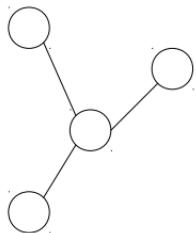


Data modeling

- Point cloud:



- Drug design:

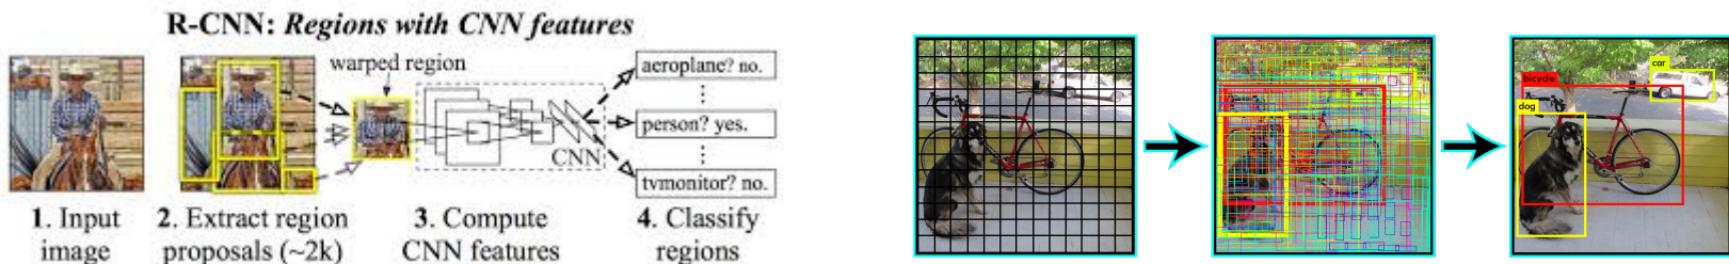


TTACGAAGGA...



Data annotation modeling

- Annotation: “ground truth” of the abstractions
- Use case specific and nontrivial...
- Example (bounding boxes): location + size + class label



R. Girshick, J. Donahue, T. Darrell, and J. Malik. Rich feature hierarchies for accurate object detection and semantic segmentation. CVPR 2014

Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi, You Only Look Once:unified, real-time object detection, CVPR 2016

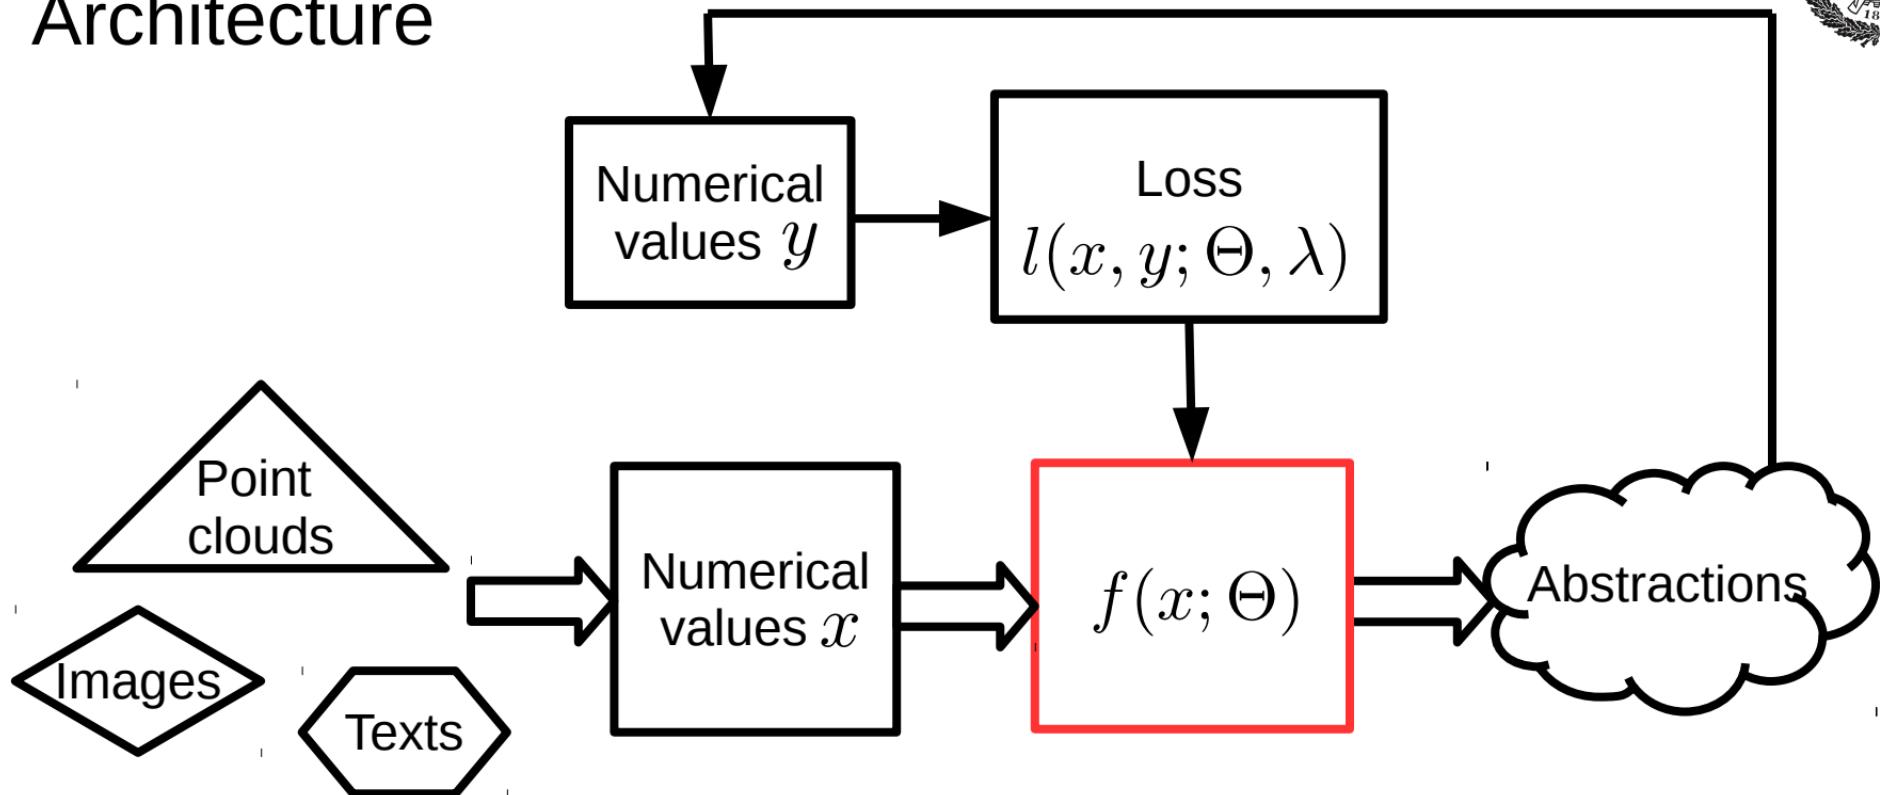


Data modeling

- Application specific
- Domain knowledge
- If you want to try out deep learning for your task
 - Be creative with data modeling
 - Apply standard network architectures



Architecture





Architecture

- Types
- Models
- Topology and wiring
- Efficiency improvement



Architecture: types

- Different structures
- Handling different inputs and producing different outputs

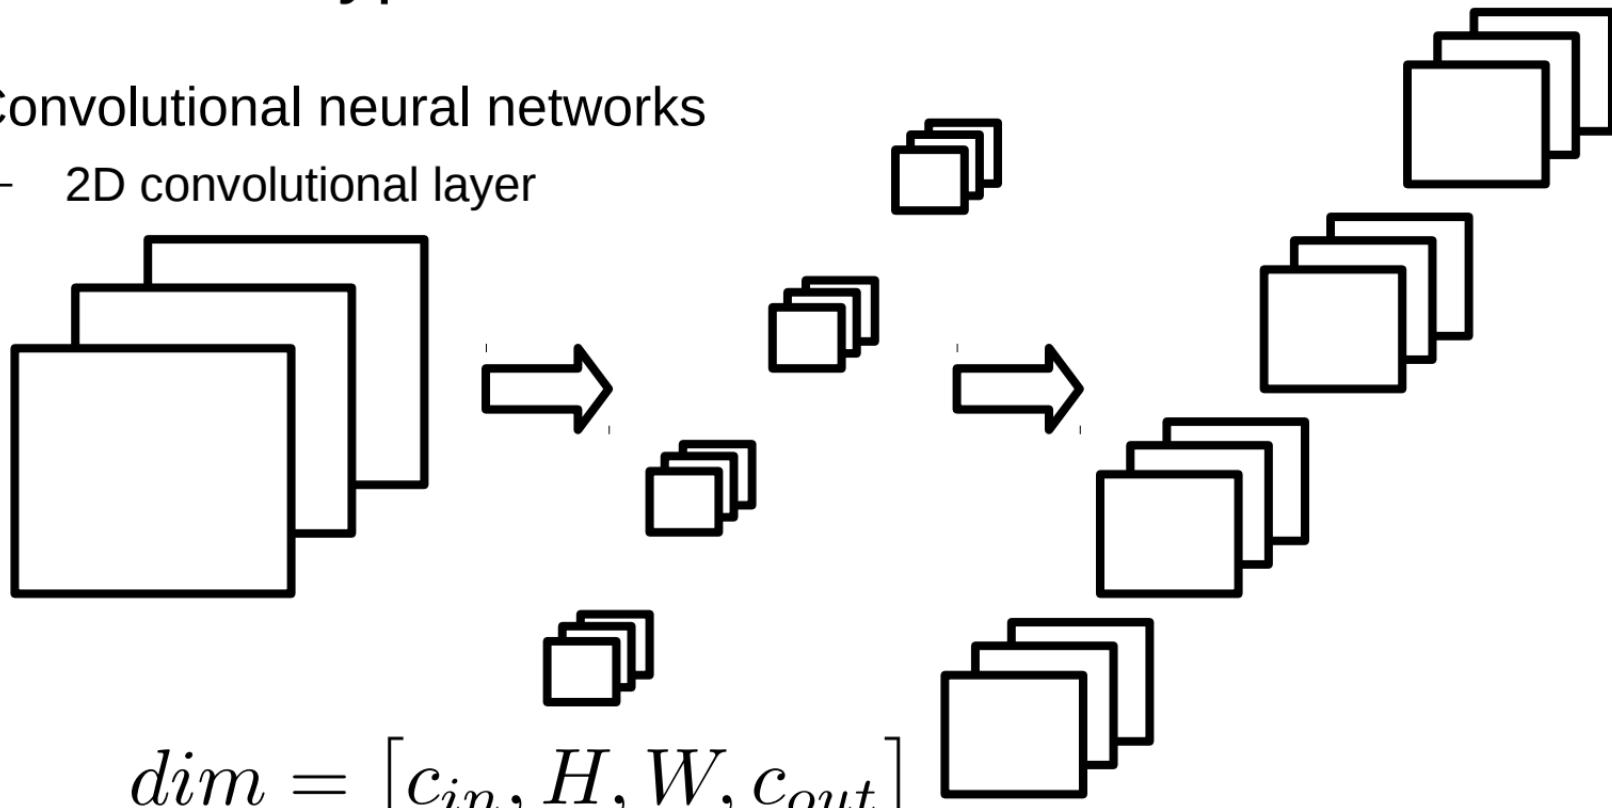


Architecture: types

- Convolutional neural networks
 - 2D convolutional layers
 - 3D convolutional layers
 - Transposed layers
 - Depth-wise separable convolutional layers
 - More convolutional layers: https://github.com/vdumoulin/conv_arithmetic
 - Pooling layers
 - Activation functions
- Artificial neural networks
- Recurrent neural networks/LSTM/neural Turing machines
- Recursive neural networks
- Graph neural networks
- Spiking neural networks

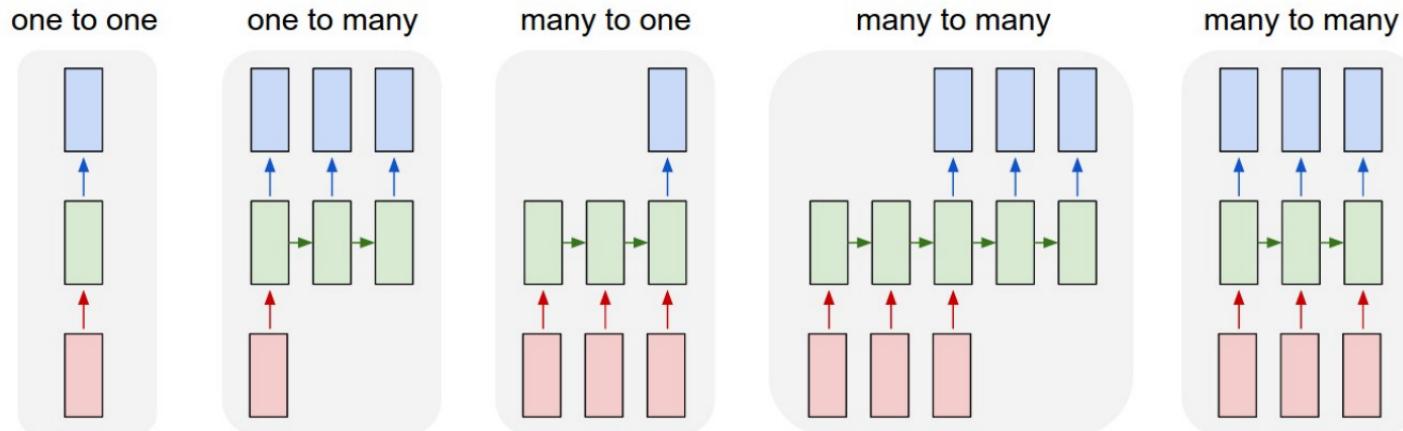
Architecture: types

- Convolutional neural networks
 - 2D convolutional layer



Architecture: types

- Recurrent neural networks





Architecture: models

- Generator/discriminator model (GAN)
- Data embedding (metric learning)
- Encoder/decoder ( )
- Attention (mainly in recurrent networks, weighted sum)



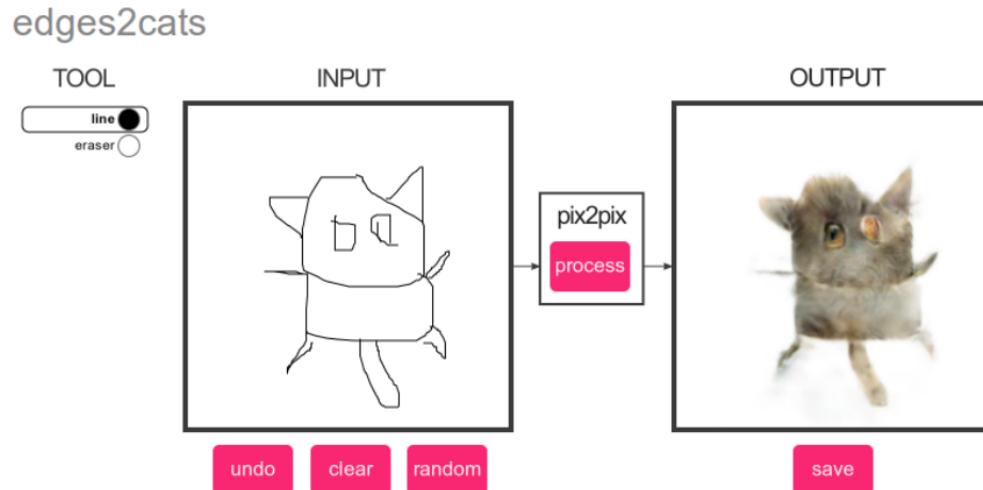
Architecture: models

- Generator/discriminator model (GAN)
- <https://affinelayer.com/pixsrv/>



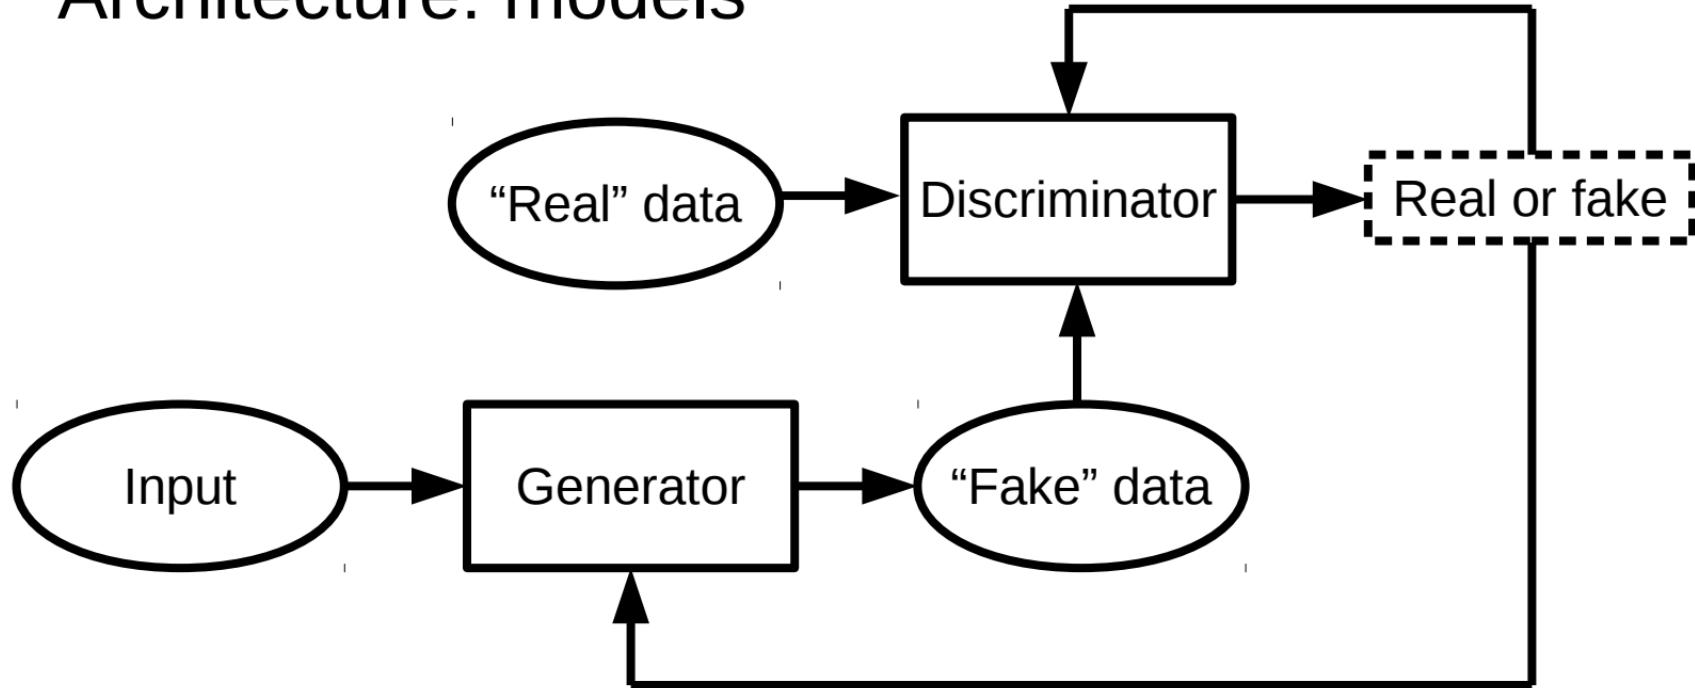
Architecture: models

- Generator/discriminator model (GAN)
- <https://affinelayer.com/pixsrv/>





Architecture: models



https://en.wikipedia.org/wiki/Generative_adversarial_network

Goodfellow Ian, Pouget-Abadie Jean, Mirza Mehdi, Xu, Bing, Warde-Farley David, Ozair Sherjil, Courville Aaron, Bengio Yoshua, Generative adversarial networks, NIPS 2014.



Architecture: models

- Generator/discriminator model (GAN)
 - “Cool” applications, “impressive” results
 - Many failure cases





Adversarial attacks



$$+ .007 \times$$



=





Adversarial training

- Deep neural nets are “too linear” → Underfitting
- Adversarial attack construction
- Defense strategies



Adversarial training

	Clean testing data	Adversarial testing data
Clean training data	★ ★	
Adversarial training data	★ ★ ★	★



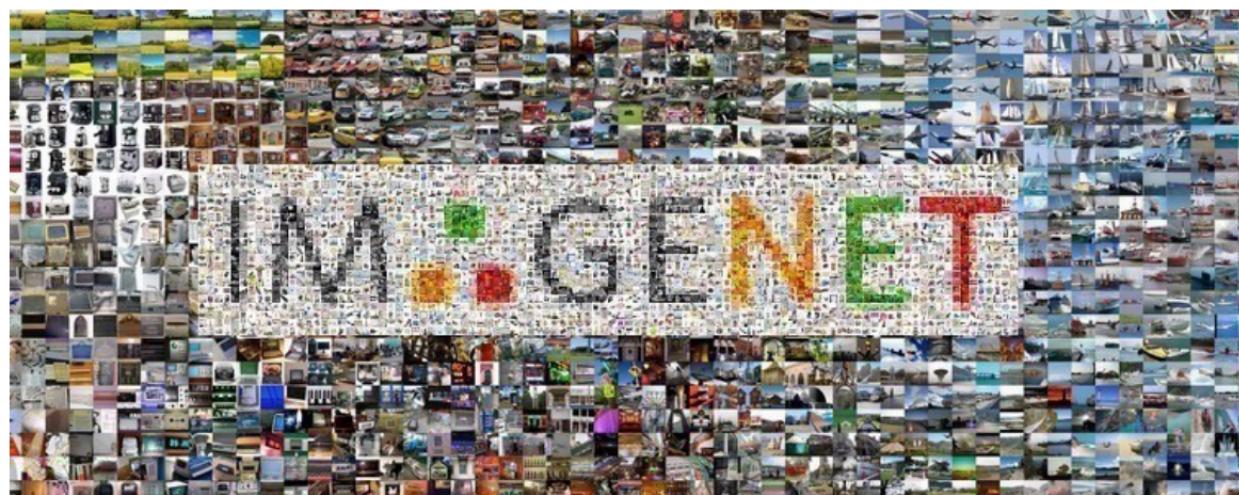
Architecture: models

- Deep embedding (metric learning)
 - Purpose:
 - find a good metric space: “similar” items should be close, “different” items should be far
 - Definition:
 - $f : \mathbb{R}^N \rightarrow \mathbb{R}^d$
 - f is a differentiable deep network
 - Different loss:
 - Margin Loss
 - Contrastive Loss
 - Triplet Loss
 - Hinge loss



Architecture: topology and wiring

- Evaluation: examples



<https://www.cs.toronto.edu/~kriz/cifar.html>
<http://www.image-net.org>



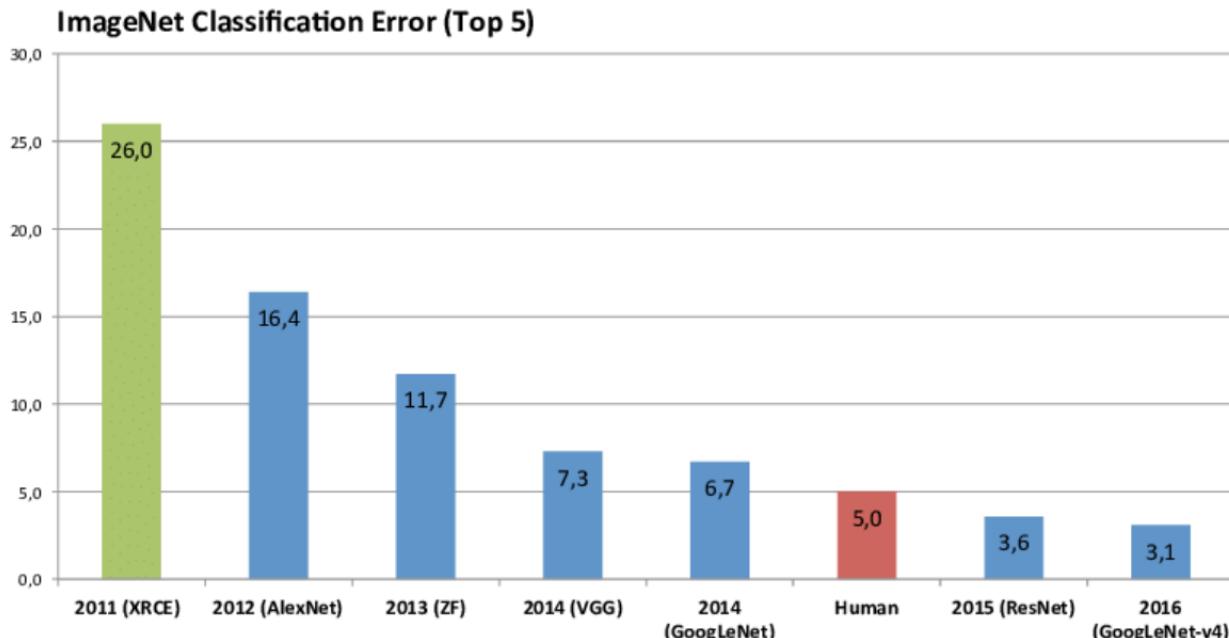
Architecture: topology and wiring

- CIFAR-10
 - 10 classes
 - 60,000 images
 - Image size: 32x32 pixels
 - Accuracy: classification rate
 - Typically takes < an hour to train
- ImageNet
 - 1000 classes
 - 1,500,000 images
 - Image size: approximately 200x200 pixels
 - Accuracy: top-1 & top-5
 - Typically takes days to train



Architecture: topology and wiring

- Evaluation result: ImageNet challenge

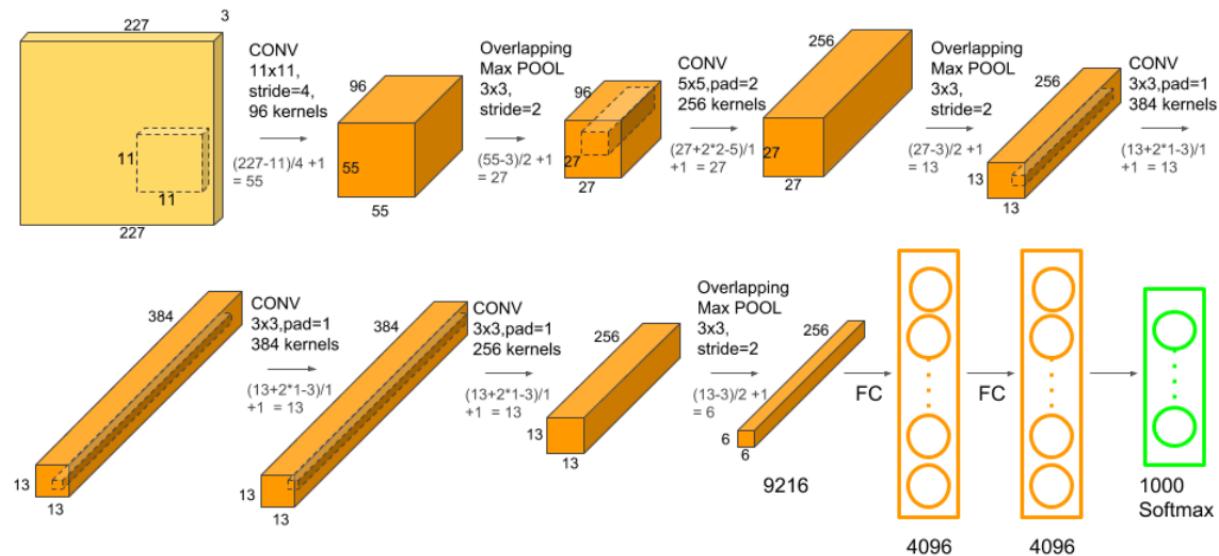




Architecture: topology and wiring

- Examples:

AlexNet



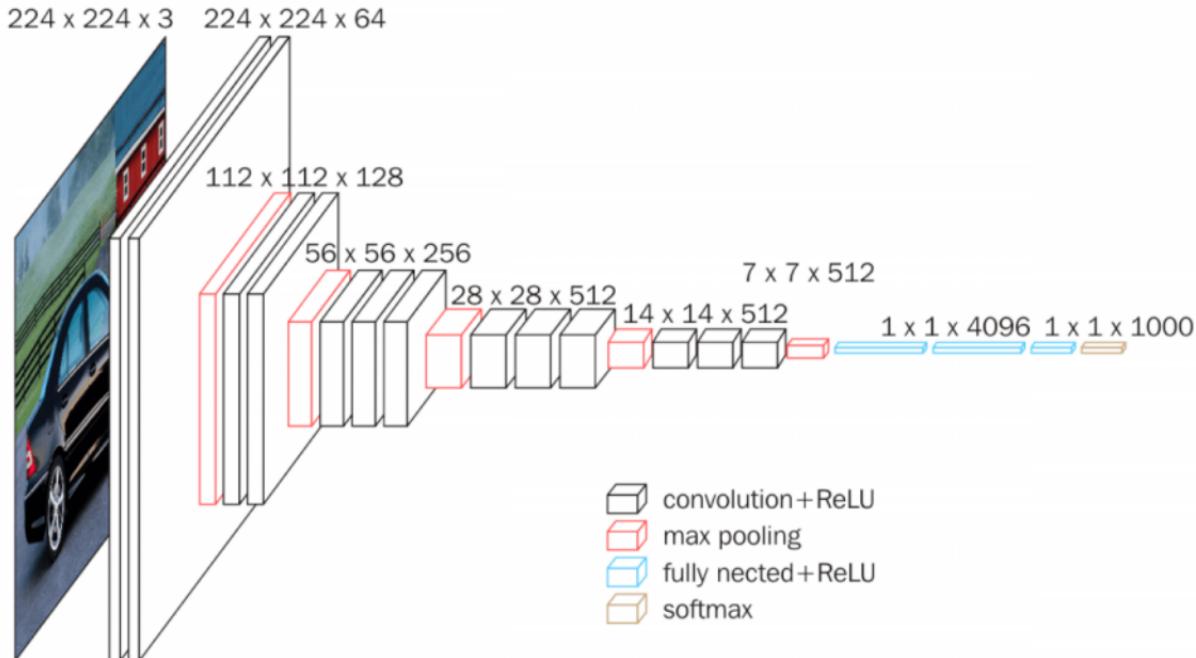
Alex Krizhevsky, ImageNet classification with deep convolutional neural networks, NIPS 2012

<https://neurohive.io/en/popular-networks/alexnet-imagenet-classification-with-deep-convolutional-neural-networks/>

Architecture: topology and wiring

- Examples:

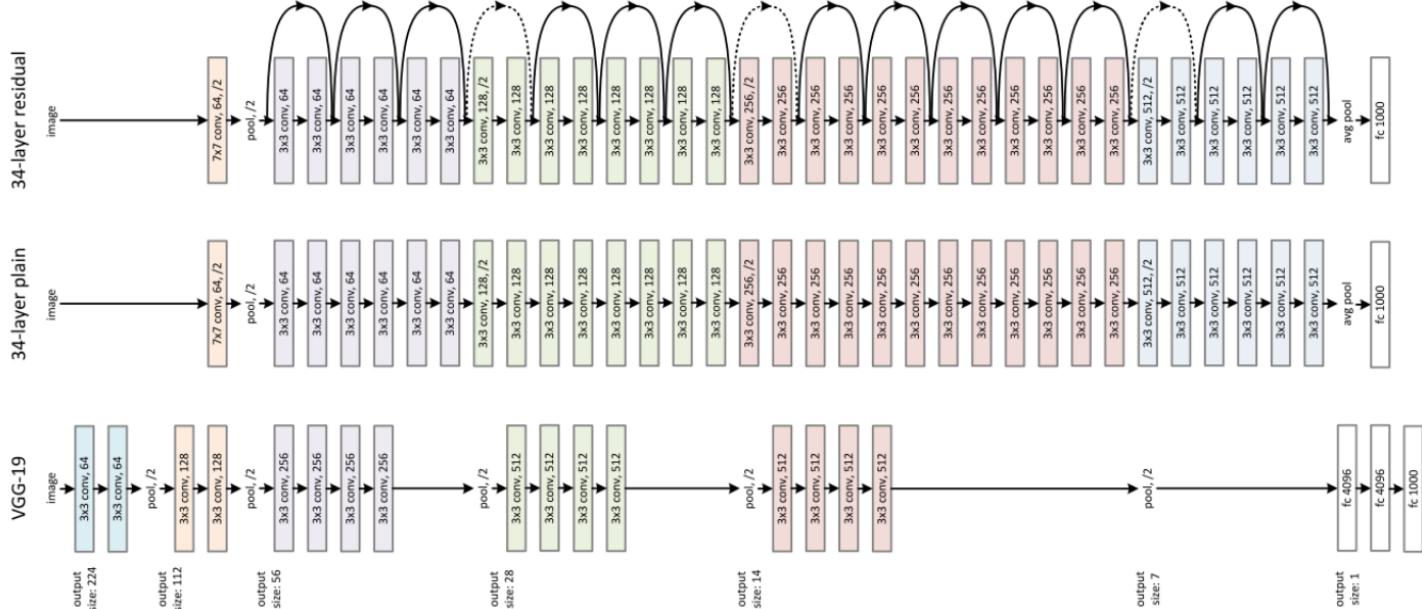
VGG



Karen Simonyan, Andrew Zisserman, Very deep convolutional networks for large-scale image recognition, ICLR 2015
<https://neurohive.io/en/popular-networks/vgg16/>

Architecture: topology and wiring

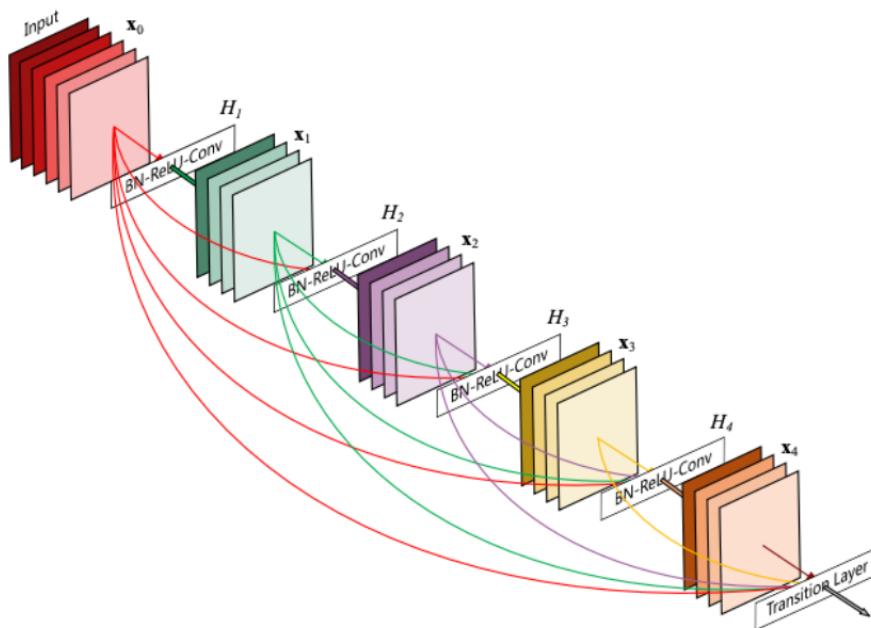
- Examples:
ResNet





Architecture: topology and wiring

- Examples:
DenseNet





Architecture: topology and wiring

Figure 1: A canonical Inception module (Inception V3).

- Examples:

Inception/Xception

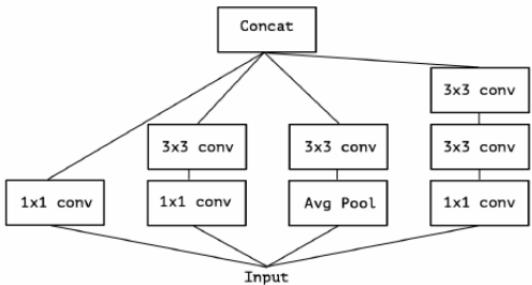


Figure 2: A simplified Inception module.

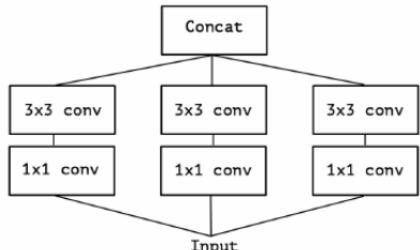
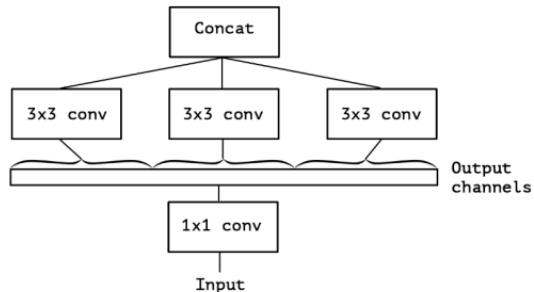
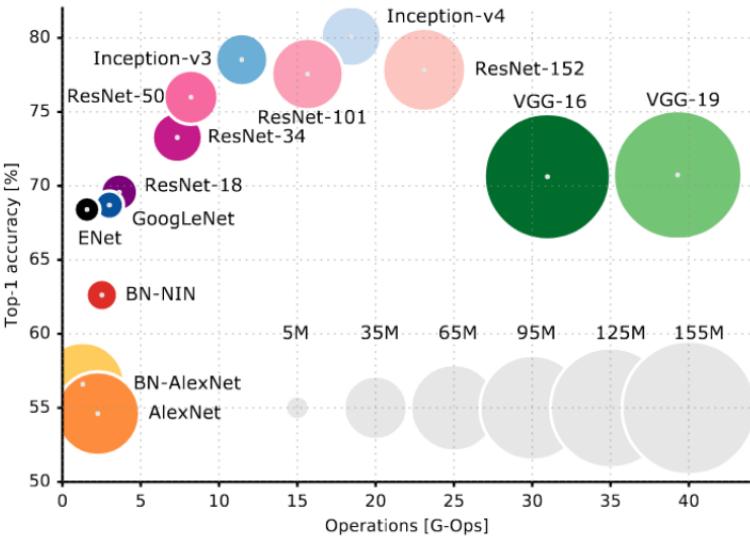
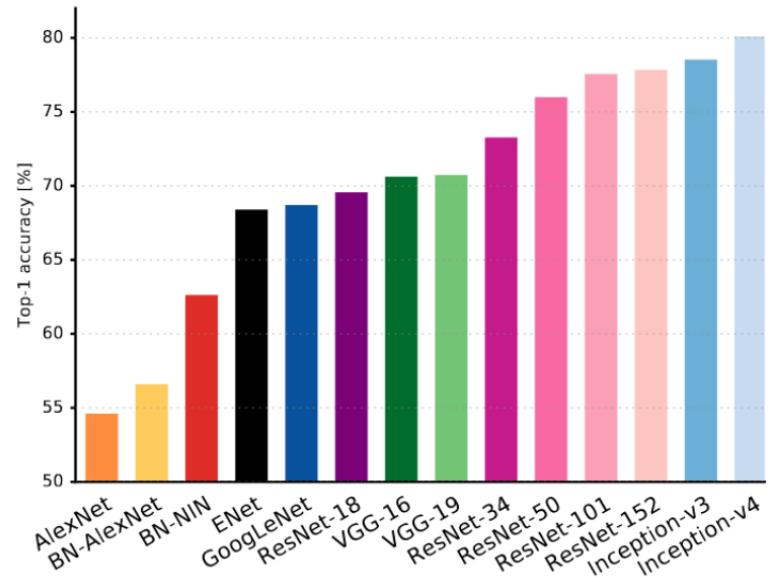


Figure 3. A strictly equivalent reformulation of the simplified Inception module.





Architecture: topology and wiring



A Canziani, A Paszke, E Culurciello, An analysis of deep neural network models for practical applications, ICLR 2017



Architecture: topology and wiring

Q: What are my choices?



Architecture: topology and wiring

Q: What are my choices?

A: Connections, filter size, number of filters, strides, activation functions, dropout, number of neurons



Architecture: topology and wiring

Q: What are my choices?

A: Connections, filter size, number of filters, strides, activation functions, dropout, number of neurons

Q: How do I know what is a good choice?



Architecture: topology and wiring

Q: What are my choices?

A: Connections, filter size, number of filters, strides, activation functions, dropout, number of neurons

Q: How do I know what is a good choice?

A: Performance on benchmark datasets



Architecture: topology and wiring

Q: What are my choices?

A: Connections, filter size, number of filters, strides, activation functions, dropout, number of neurons

Q: How do I know what is a good choice?

A: Performance on benchmark datasets

Q: Why did I choose that architecture?



Architecture: topology and wiring

Q: What are my choices?

A: Connections, filter size, number of filters, strides, activation functions, dropout, number of neurons

Q: How do I know what is a good choice?

A: Performance on benchmark datasets

Q: Why did I choose that architecture?

A: $\neg(\psi) \vdash$



Architecture: topology and wiring

- Network Architecture Search (NAS)
 - Reinforcement learning
 - Progressive network search
 - Differentiable architecture search
- Macro vs micro search
 - Macro search: “give me the best network”
 - Micro search: “given a skeleton of the network, give me the best wiring/components for each block”

https://en.wikipedia.org/wiki/Neural_architecture_search

<https://www.automl.org/automl/literature-on-neural-architecture-search/>

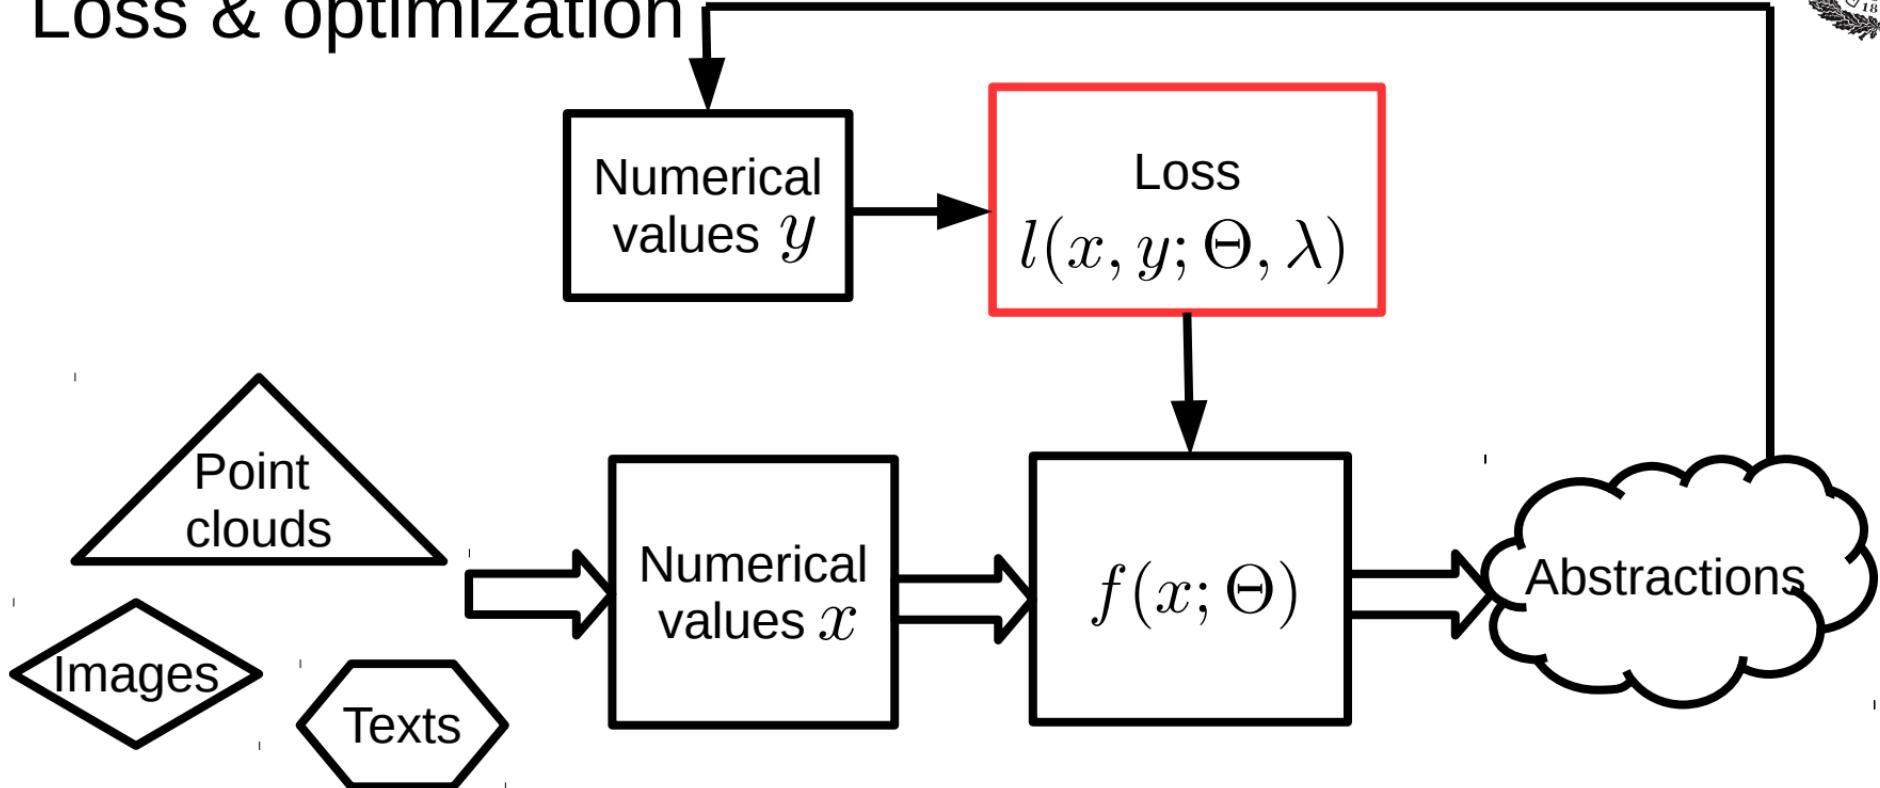


Architecture: efficiency improvement

- Pruning (remove insignificant weights/connections)
- Compression (quantization, Huffman encoding, binarization)
- Weight sharing (use the same weights to reduce the dof)
- Low rank approximation (tensor/matrix factorization)



Loss & optimization





Loss functions

- Cross entropy loss: $l_i = -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$
- Hinge loss: $l_i = \sum_{j \neq y_i} \max(0, \hat{y}_j - \hat{y}_{y_i} + 1)$
- Mean Absolute Error (L1 Loss): $l_i = |\hat{y}_i - y_i|$
- Mean Square Error/Quadratic Loss (L2 Loss): $l_i = (\hat{y}_i - y_i)^2$

i : data point index

y : label

\hat{y} : prediction from the network

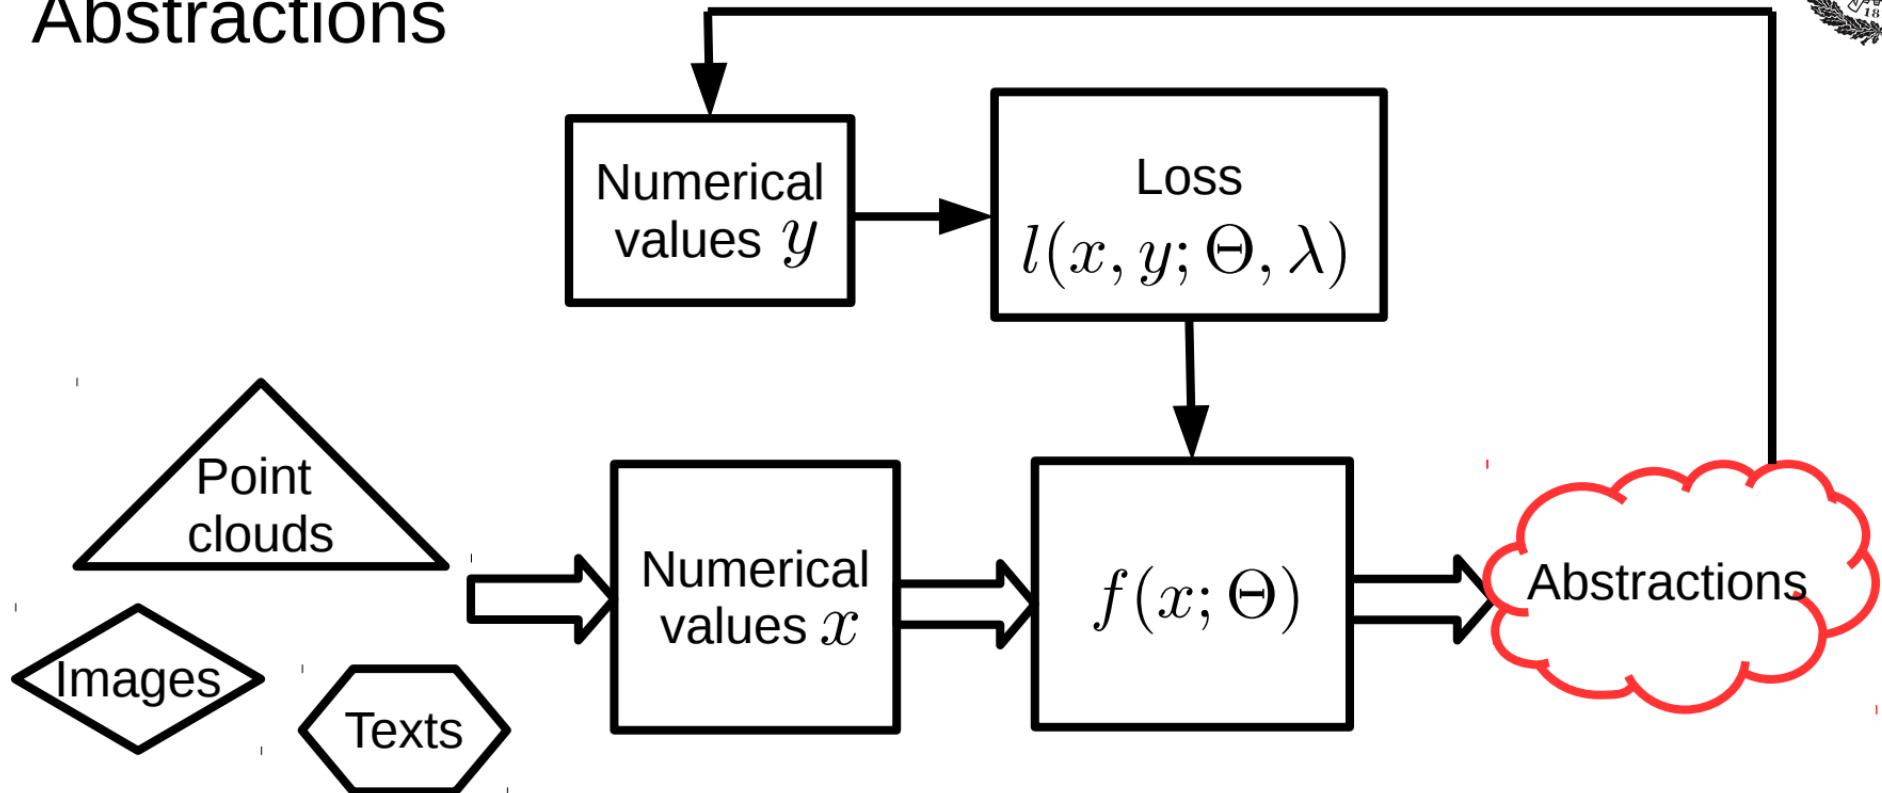


Optimization

- Optimizer
 - Stochastic Gradient Descent
 - Momentum (“moving average”)
 - ADAM
- Learning rate decay (mainly for first order optimization)
- Regularization (“specificity” → “generalization”)
 - Empirical risk minimization & statistical learning theory
 - Drop out (set activation to zero)
 - Batch normalization
 - Data augmentation
 - Ensemble learning



Abstractions



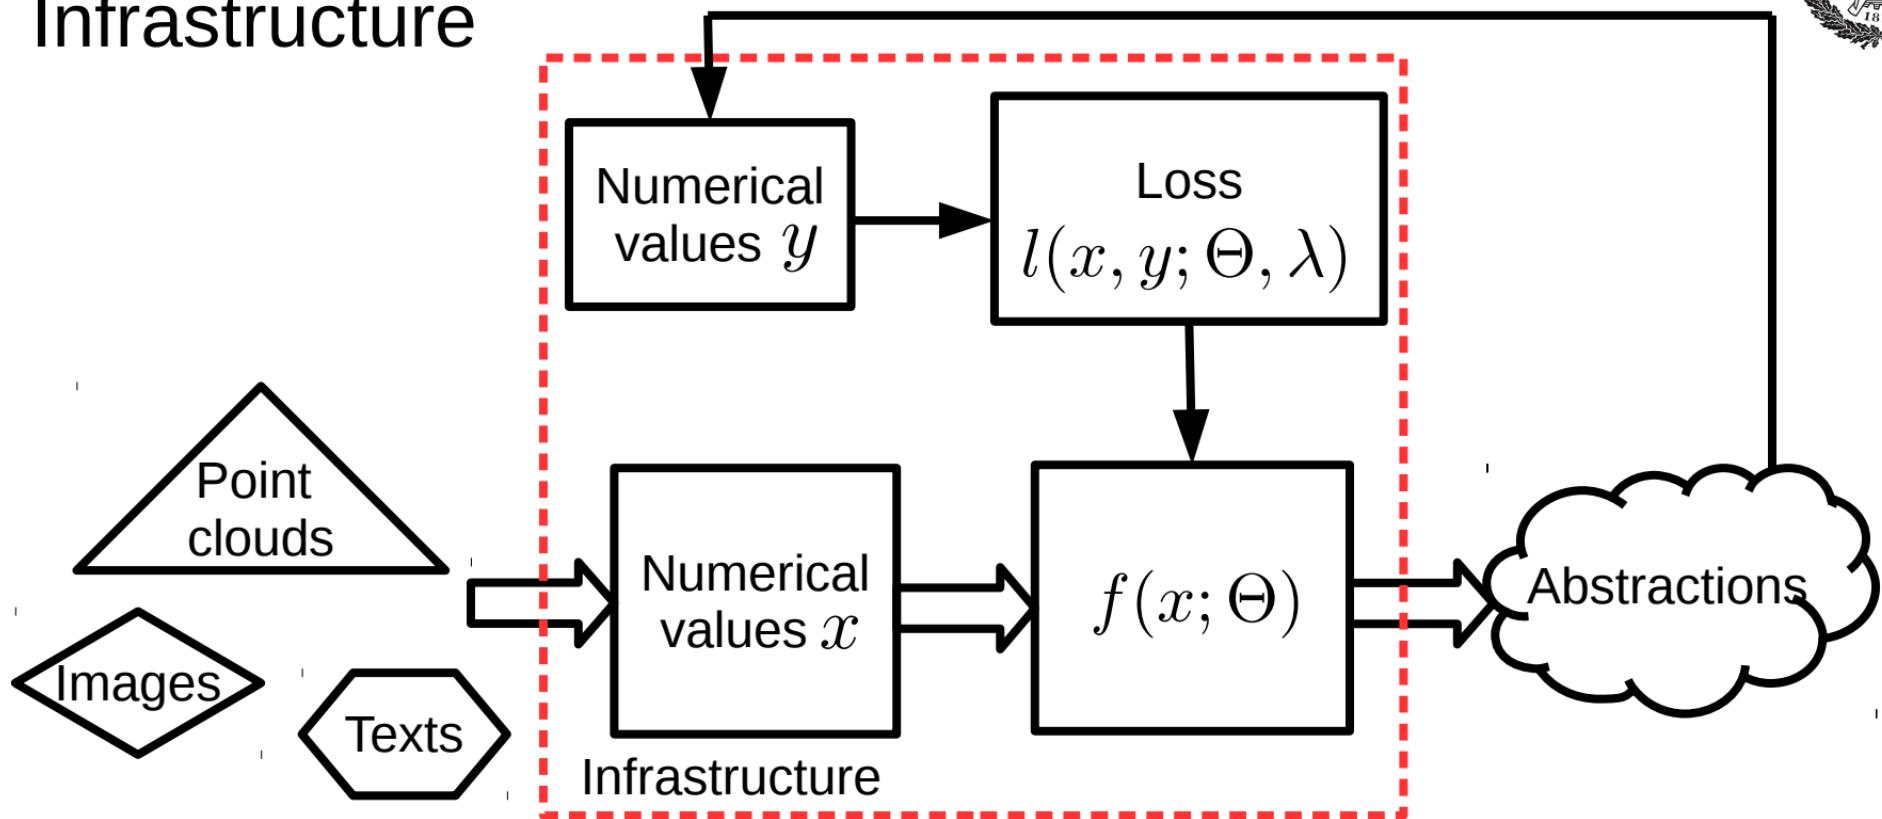


Abstractions

- Annotations for supervised learning
 - Annotation guideline
 - Manual annotation
 - Automation
 - Clustering, semi-supervised learning, crowd sourcing and queries
 - KPI of use cases



Infrastructure





Infrastructure

- CPU/GPU/storage
- CPU: preparing data
- GPU: backprops and inference
 - Laptop/workstation
 - GPU servers
- Storage: fast random access
- Hosting/cooling



Conclusion

- Deep learning works
- Deep learning evolves fast
- Deep learning brings everything together
 - “Big data” is here
 - Modern hardware + low level libraries enable scalability
 - Data modeling for fancy applications
 - Highly complex models: low bias, high variance
 - Regularization to balance bias/variance trade-off
 - Optimizers: faster and better
- Deep learning is mysterious
 - Adversarial attacks
 - Automation, analysis and verification

MLFP



Thank you